

Normalization in LLMs

- Normalization in LLMs is typically applied between the attention and Feed forward layers (and vice versa) it's usually applied before or after each sub layer (depending on variants):
 - post norm: normalization after Attention/MLP (used in original transformer paper)
 - pre norm: normalization before Attention/MLP (used in modern LLMs as it stabilizes training in deep models)
- Why?:
 - prevents exploding/vanishing activations: means the numbers (output of layers) don't become too big or small as they pass through many layers - this keeps the model's computation stable
 - makes optimizations smoother: means the model learns more steadily during training, the loss surface becomes less chaotic, so gradient descent can take more reliable steps towards better performance instead of bouncing around / getting stuck
 - allows deeper, more stable transformer training: because norm keeps each layers well behaved (balanced and stable not too big or small, consistent across layers) you can stack more layers without the model's training becoming unstable or diverging. In short normalization makes it possible for very deep NN. like LLMs to still learn effectively and converge smoothly, instead of breaking down as they grow in size